RAGAS Evaluation Report

Retrieval-Augmented Generation System

Evaluation Date: 2026-09-01T15:03:44.531313
Total Questions: 10

OVERALL RESULT: FAIL

Evaluation Summary

Metric	Score	Required	Status
Faithfulness	77.50%	> 90%	FAIL
Answer Correctness	73.00%	> 80%	FAIL
Context Recall	79.00%	> 85%	FAIL
Context Precision	78.50%	> 80%	FAIL

Required Minimum Scores

Metric	Minimum
Faithfulness	> 90%
Answer Correctness	> 80%
Context Recall	> 85%
Context Precision	> 80%

Question-Level Evaluation

Question	Faith.	Correct.	Recall	Precision
Q1: What problem does the Transformer architecture address, and how does it differ from recurrent and convolutional sequence models?	0.0%	0.0%	0.0%	0.0%
Q2: What is scaled dot-product attention in the Transformer, and why is the scaling factor 1/sqrt(d_k) used?	0.0%	0.0%	0.0%	0.0%
Q3: What is BERT, and what is the main advantage of its bidirectional pre-training approach?	0.0%	0.0%	0.0%	0.0%
Q4: What are the two pre-training tasks used to train BERT?	0.0%	0.0%	0.0%	0.0%
Q5: What is RoBERTa, and what did the authors conclude about the original BERT pre-training procedure?	0.0%	0.0%	0.0%	0.0%
Q6: What major training changes did RoBERTa make compared with BERT to improve performance?	0.0%	0.0%	0.0%	0.0%
Q7: What is the main idea behind the T5 text-to-text framework, and how does it unify different NLP tasks?	0.0%	0.0%	0.0%	0.0%
Q8: What is the Colossal Clean Crawled Corpus (C4), and why was it important for the T5 study?	0.0%	0.0%	0.0%	0.0%
Q9: What is GPT-3's few-shot learning approach, and how does it differ from traditional fine-tuning?	0.0%	0.0%	0.0%	0.0%
Q10: What were the main capabilities and limitations observed for GPT-3 in the few-shot evaluation?	0.0%	0.0%	0.0%	0.0%